

[< back](#) [next >](#)

VIP



Retrieve data using an index #NZ8 Pending

Complete previous stages to gain access to this stage.



Instructions

[Code Examples](#)[Forum](#)

Your Task Pending

HARD

In this stage, we'll implement an index scan. Rather than reading *all* rows in a table and then filtering in-memory, we'll use an index to perform a more intelligent search.

To test whether your implementation actually uses an index, the tester will use a database is ~1GB in size and expect your program to return query results in less than 3 seconds.

The test database contains a `companies` table with an index named `idx_companies_country` on the `country` column.

You can download a small version of this database to test locally, read the **Sample Databases** section in the **README** of your repository for details.

Here's how the tester will execute your program:

```
$ ./your_sqlite3.sh companies.db "SELECT id, name FROM companies WHERE count
```

and here's the output it expects:

```
121311|unilink s.c.  
2102438|orange asmara it solutions  
5729848|zara mining share company
```

Ready to run tests...

[Show logs](#)